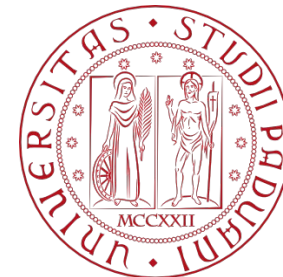


# COMPUTER ENGINEERING LABORATORY

**Luigi Rizzo**

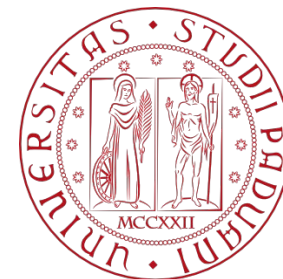
[luigi.rizzo@unipd.it](mailto:luigi.rizzo@unipd.it)

October 2024-January 2025



UNIVERSITÀ  
DEGLI STUDI  
DI PADOVA

# Socket programming



UNIVERSITÀ  
DEGLI STUDI  
DI PADOVA

# The client-server model



In modern operating systems, the services available on the network are mainly based on the client/server model. This architecture allows systems to share resources and cooperate to achieve an objective through the presence of two categories of subjects, service programs, called servers, which receive requests and provide responses, and user programs, called clients.


A server is (normally) able to respond to more than one client, so it is possible that many programs can interact simultaneously. What distinguishes the model, however, is that the architecture of the interaction is always in terms of many towards one, the server, who comes to take on a privileged role.

All the fundamental services of the Internet follow this model, such as web pages, e-mail, ftp, telnet, and practically every service that is provided over the network.

# The client-server model




- ❑ Computer Network
    - ❑ hosts, routers, communication channels
  - **Hosts** run applications
  - **Routers** forward information
  - **Packets**: sequence of bytes
    - contain control information
    - e.g. destination host
  - **Protocol** is an agreement
    - meaning of packets
    - structure and size of packets
- e.g. Hypertext Transfer Protocol (HTTP)




- Several protocols for different problems
  - **Protocol Suites or Protocol Families:** TCP/IP
- TCP/IP provides end-to-end connectivity specifying how data should be
  - formatted,
  - addressed,
  - transmitted,
  - routed, and
  - received at the destination
- can be used in the internet and in stand-alone private networks
- it is organized into **layers**

# Protocol Families -TCP/IP

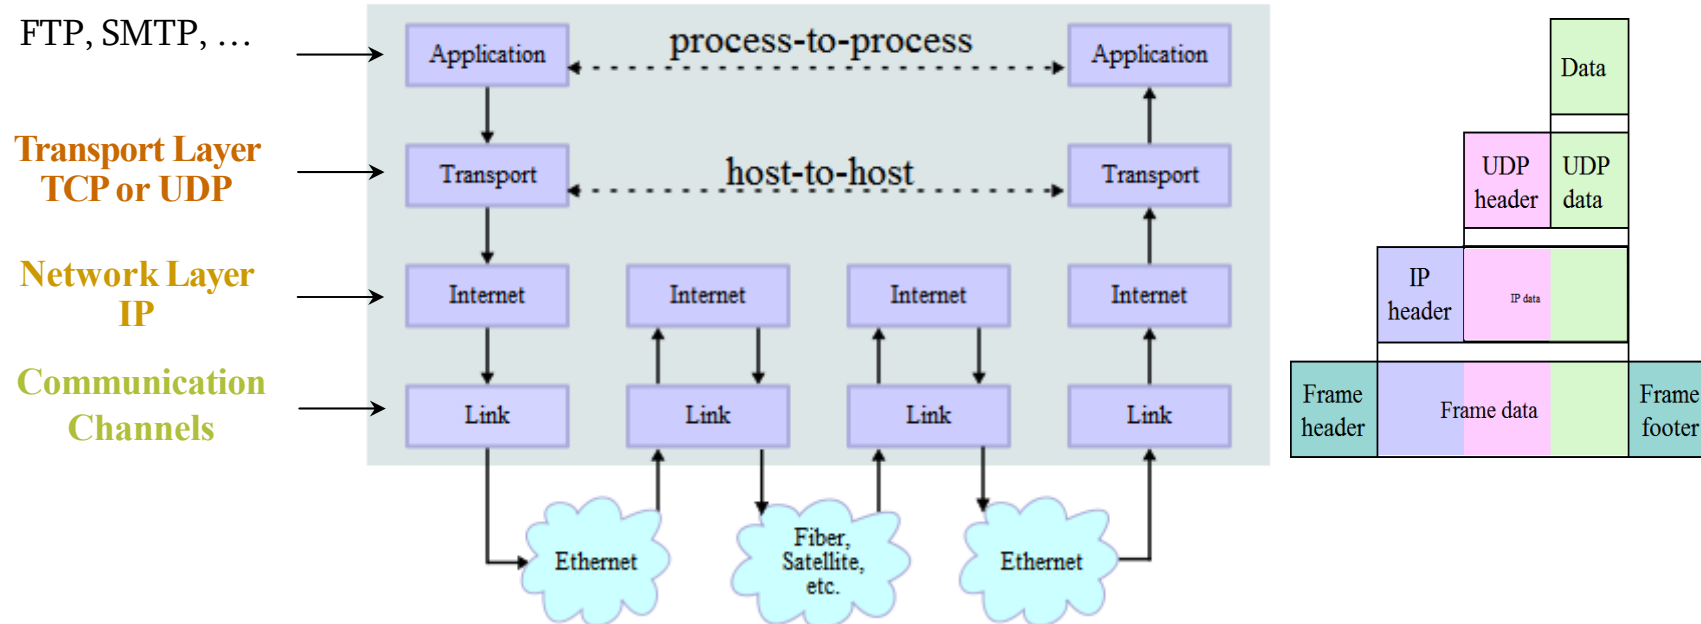


## Network Topology



\*

## Data Flow



\* image is taken from [http://en.wikipedia.org/wiki/TCP/IP\\_model](http://en.wikipedia.org/wiki/TCP/IP_model)

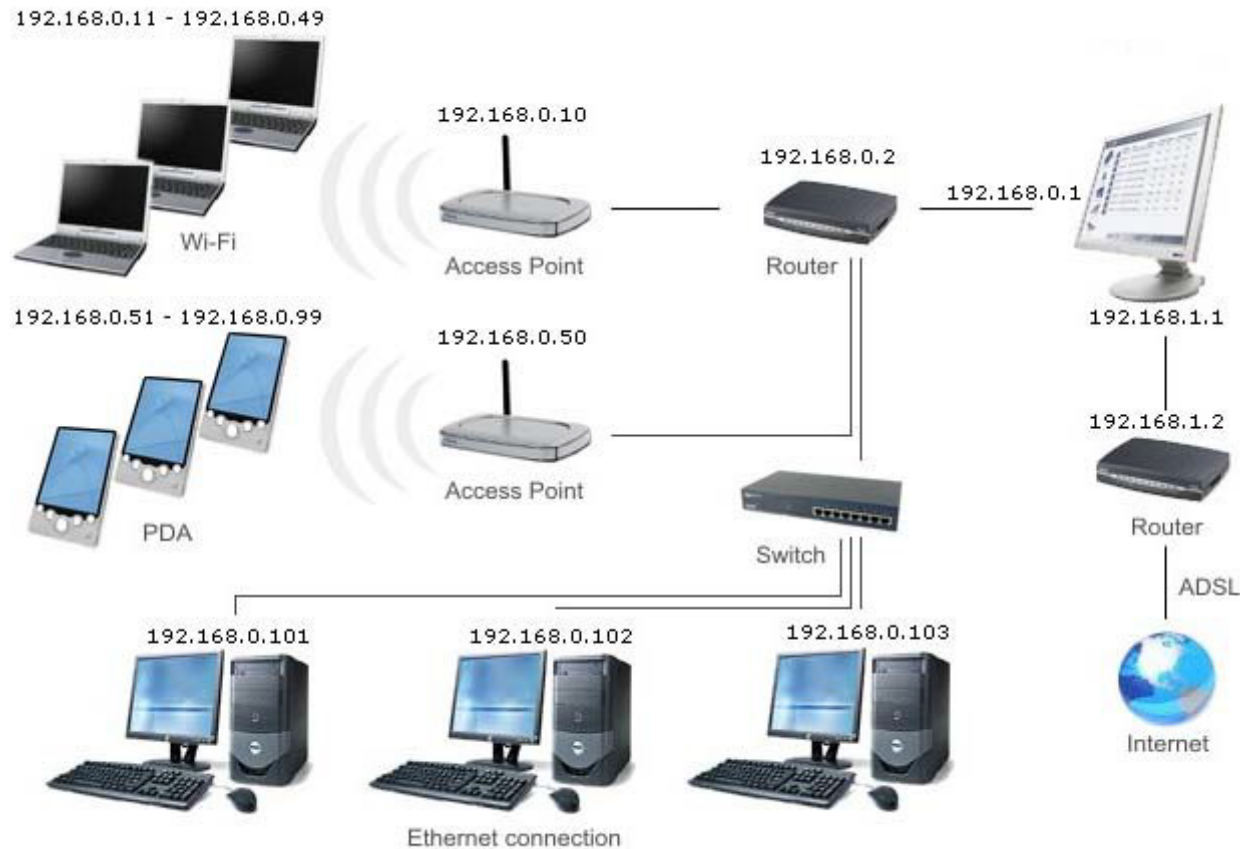
Concerning system programming, there is the socket interface which provides a user-friendly abstraction of the basic mechanisms for implementing client/server programs. A socket (“**socket**”) is a communication end between processes.

Socket Programming is a method to connect two nodes over a network to establish a means of communication between those two nodes. A node represents a computer or a physical device with an internet connection. A socket is the endpoint used for connecting to a node. The signals required to implement the connection between two nodes are sent and received using the sockets on each node respectively.

# Local Area Network Addresses - IPv4



- The 32 bits of an IPv4 address are broken into 4 octets, or 8 bit fields (0-255 value in decimal notation).





- Both use **port numbers**
  - application-specific construct serving as a communication endpoint
  - 16-bit unsigned integer, thus ranging from 0 to 65535
    - to provide **end-to-end** transport
- UDP: User Datagram Protocol
  - no acknowledgements
  - no retransmissions
  - out of order, duplicates possible
  - connectionless, i.e., app indicates destination for each packet
- TCP: Transmission Control Protocol
  - reliable **byte-stream channel** (in order, all arrive, no duplicates)
    - similar to file I/O
  - flow control
  - connection-oriented
  - bidirectional

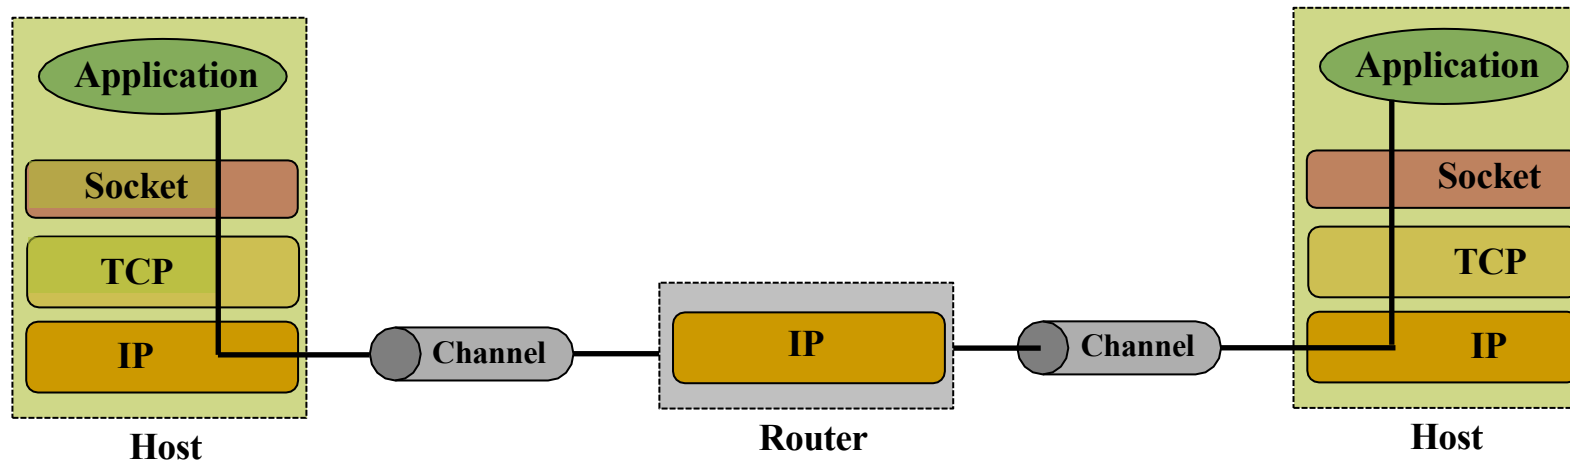
- TCP is used for services with a large data capacity, and a persistent connection
- UDP is more commonly used for quick lookups, and single use query-reply actions.
- Some common examples of TCP and UDP with their default ports:

DNS lookup	UDP	53	Domain Name System lookup
FTP	TCP	21	
HTTP	TCP	80	
SNMP	UDP	161	Simple Network Management Protocol
Telnet	TCP	23	
DHCP	UDP	67/68	Dynamic Host Configuration Protocol

# Berkley Sockets



- Universally known as **Sockets**
- It is an abstraction through which an application may send and receive data
- Provide **generic access** to interprocess communication services
- Standard API for networking



- ❑ Uniquely identified by
  - an internet address
  - an end-to-end protocol (e.g. TCP or UDP)
  - a port number
- ❑ Two types of (TCP/IP) sockets
  - Stream sockets (e.g. uses TCP)
    - provide reliable byte-stream service
  - Datagram sockets (e.g. uses UDP)
    - provide best-effort datagram service
    - messages up to 65.500 bytes
- ❑ Socket extend the convectional UNIX I/O facilities
  - file descriptors for network communication
  - the read and write system calls are extended

internal data structure

Family: AF_INET
Service: SOCK_STREAM
Local_IP:
Remote_IP:
Local_Port:
Remote_Port:
...

In simpler terms, a socket allows applications to talk to each other, either over the internet (e.g., between a client and server) or within the same computer.

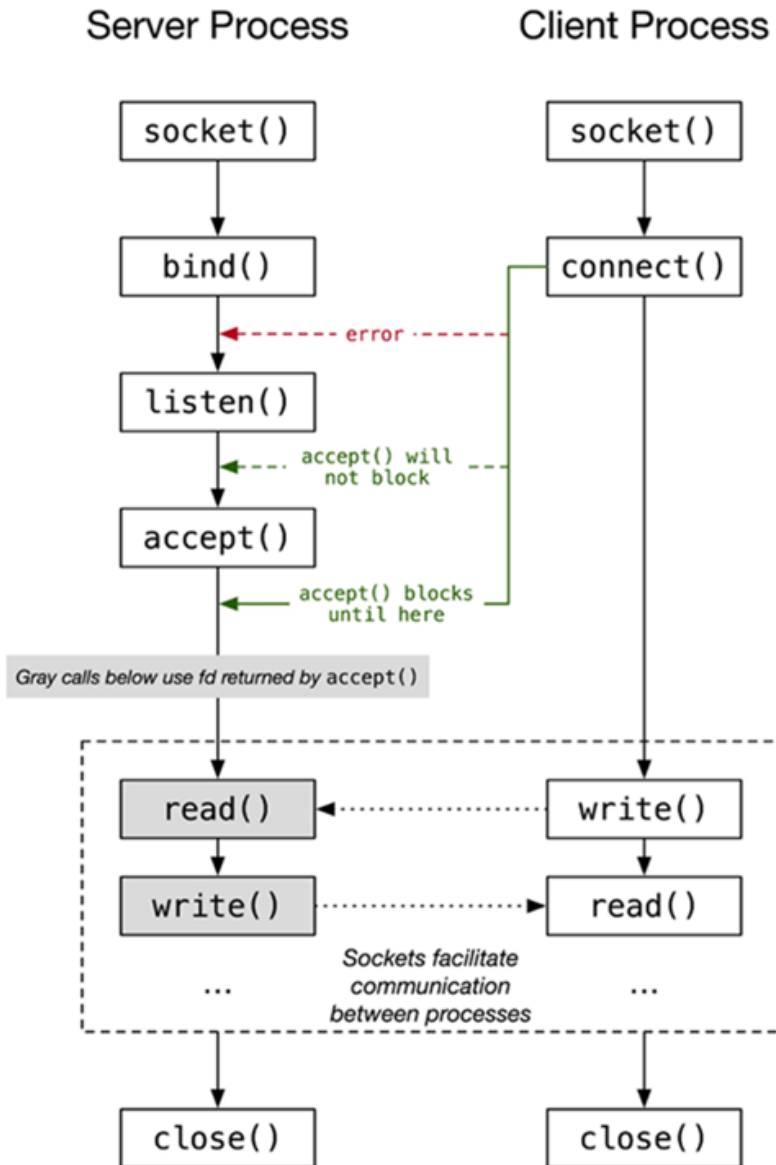
# The client-server model



Servers are normally divided into two main categories, and are called concurrent or iterative, based on their behavior.

- An iterative server responds to the request by sending data and remains busy and does not respond to further requests until it has provided a response to the request. Once the response is complete the server becomes available again.
- A concurrent server, instead, when processing the request, creates a child process (or a thread) responsible for providing the requested services, to immediately wait for further requests. In this way, with multitasking systems, multiple requests can be satisfied simultaneously. Once the child process has finished its work it is usually terminated, while the original server always remains active.

# Client and server model state diagram



- The nodes are divided into two types, server node and client node.
- The client node sends the connection signal and the server node receives the connection signal sent by the client node.
- The connection between a server and client node is established using the socket over the transport layer of the internet.
- After a connection has been established, the client and server nodes can share information between them using the read and write commands.
- After sharing of information is done, the nodes terminate the connection.

Different stages must be performed on the server node to receive a connection sent by the client node.

- Socket creation
- Setsockopt
- Bind
- Listen
- Accept
- Send and receive data. There are a number of ways to do this, but the simplest is to use the `read()` and `write()` / `send()` system calls.

The client-side sends the connection requests to the server-side. To send these requests several stages have to be performed on the client side too.

- Socket Connection
- Connect
- Send and receive data. There are a number of ways to do this, but the simplest is to use the `read()` and `write()` / `send()` system calls.



- Socket creation

The first stage deals with the creation of a socket, which is the basic component for sending or receiving signals between nodes. The `sys/socket.h` header has the necessary functions to create a socket in C. In socket programming in C, a socket can be created by the `socket()` function with syntax,

```
int socket(int domain, int type, int protocol);
```

example

```
int server_fd = socket(AF_INET, SOCK_STREAM, 0);
```

*int socket(int domain, int type, int protocol);*

The **domain** represents the address family over which the communication will be performed. The domain is pre-fixed values present in the sys/socket.h header.

Some domains are,

- **AF\_LOCAL** or **AF\_UNIX** is used for local communication or in the case where the client and server are on the same node. These sockets are called **UNIX domain sockets**.
- **AF\_INET** is used to represent the IPv4 address of the client to which a connection should be made. Similarly, **AF\_INET6** is used for IPv6 addresses. These sockets are called **internet domain sockets**.
- **AF\_BLUETOOTH** is used for low-level Bluetooth connection.

*int socket(int domain, int type, int protocol);*

The **type** represents the type of communication used in the socket. Some mostly used types of communication are,

- **SOCK\_STREAM** uses the TCP (Transmission Control Protocol) to establish a connection. This type provides a reliable byte stream of data flow and is a connection-based protocol. These sockets are called **stream sockets**.
- **SOCK\_DGRAM** uses the UDP (User Datagram Protocol) which is unreliable and a connectionless protocol. These sockets are also called **datagram sockets**.

*int socket(int domain, int type, int protocol);*

The **protocol** represents the protocol used in the socket. This is represented by a number. When there is only one protocol in the protocol family, the protocol number will be 0, or else the specific number for the protocol has to be specified.

The **socket()** function creates a socket and returns a file descriptor which represents an open file that will be utilized by the socket in reading and writing operations and the file descriptor is used to represent the socket in later stages. In case of an error in creating the socket, -1 is returned by the socket() function.

- Setsockopt

The `setsockopt()` function in socket programming in C is used to specify some options for the socket to control the behavior of the socket. The syntax is,

```
int setsockopt(int socket_descriptor, int level, int option_name, const void  
*value_of_option, socklen_t option_length);
```

example

```
int opt = 1;  
setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR, &opt, sizeof(opt));
```

```
int setsockopt(int socket_descriptor, int level, int option_name, const void  
*value_of_option, socklen_t option_length);
```

The **socket\_descriptor** is the value of the file descriptor returned by the `socket()` function.

The **level** parameter represents the level at which the option for the socket must be applied. The `SOL_SOCKET` represents the socket level and `IPPROTO_TCP` represents the TCP level.

*int setsockopt(int socket\_descriptor, int level, int option\_name, const void \*value\_of\_option, socklen\_t option\_length);*

The **option\_name** specifies the rules or options that should be modified for the socket. Some useful options are,

- SO\_DEBUG is used to enable the recording of debugging information.
- SO\_REUSEADDR is used to enable the reusing of local addresses in the bind() function.
- SO\_SNDBUF is used to set the maximum buffer size that can be sent using the socket connection.
- SO\_LINGER is used to set that socket lingers on close.

```
int setsockopt(int socket_descriptor, int level, int option_name, const void  
*value_of_option, socklen_t option_length);
```

The **option\_value** is used to specify the value for the options set in the option\_name parameter.

The **option\_length** is the length of the variable used to set the option value.

The function returns a value of 0 of data type int on the successful application of the option and a value of -1 on failure.



- Bind

The `bind()` function in socket programming in C is used to assign an address to a socket created using the `socket()` function. The syntax of `bind()` function is,

```
int bind(int socket_descriptor , const struct sockaddr *address, socklen_t  
length_of_address);
```

The **socket\_descriptor** is the value of the file descriptor returned by the `socket()` function.

```
int bind(int socket_descriptor , const struct sockaddr *address, socklen_t  
length_of_address);
```

The **address** is a structure of type `sockaddr`. We usually use a structure of type `sockaddr_in` to represent this information, because information such as port and address can only be stored in this structure. The `sockaddr_in` is cast to the `sockaddr` data type when calling the `bind()` function.

The **length\_of\_address** represents the size of the address passed as the second parameter.

The function returns 0 on binding the address and port successfully or returns -1 on failure.

```
int bind(int socket_descriptor , const struct sockaddr *address, socklen_t  
length_of_address);
```

example

```
struct sockaddr_in address;  
socklen_t addrlen = sizeof(address);
```

```
address.sin_family = AF_INET;  
address.sin_addr.s_addr = INADDR_ANY;  
address.sin_port = htons(PORT);  
bind(server_fd, (struct sockaddr*)&address, sizeof(address));
```

# struct sockaddr\_in



The struct **sockaddr\_in** data type is a specialized structure used to represent IPv4 socket addresses in socket programming. It is built upon the generic struct `sockaddr` and includes fields specific to IPv4 addresses, such as the IP address and port number.

```
struct sockaddr_in {  
    sa_family_t  sin_family; // Address family (e.g., AF_INET)  
    uint16_t     sin_port;   // Port number (network byte order)  
    struct in_addr sin_addr;  // IPv4 address (32-bit)  
    char         sin_zero[8]; // Padding to match size of sockaddr  
};
```

## Fields in struct sockaddr\_in

- `sin_family`:
  - specifies the address family.
  - for IPv4, this is always set to `AF_INET` (Address Family - Internet).

*`address.sin_family = AF_INET;`*

- `sin_port`:
  - represents the port number on which the socket communicates.
  - it must be stored in network byte order (big-endian format).
  - use the `htons()` function (host-to-network short) to convert the port number.

*`address.sin_port = htons(8080); // Convert 8080 to network byte order`*

- `sin_addr`:
  - represents the IPv4 address.
  - it is defined as:

```
struct in_addr {  
    uint32_t s_addr; // 32-bit IPv4 address (network byte order)  
};
```
  - Use the `inet_addr()` or `inet_pton()` functions to convert a string IP address to the required format.

`address.sin_addr.s_addr = inet_addr("127.0.0.1");`

or

`inet_pton(AF_INET, "127.0.0.1", &address.sin_addr);`

- `sin_zero`:
  - an unused field consisting of 8 bytes of padding.
  - it exists to make `sockaddr_in` the same size as the generic struct `sockaddr`.
  - it is typically set to 0 using `memset()`:  
*`memset(address.sin_zero, 0, sizeof(address.sin_zero));`*

The struct `sockaddr` is a generic socket address structure in C typically defined as:

```
struct sockaddr {  
    sa_family_t sa_family; // Address family (e.g., AF_INET for IPv4)  
    char sa_data[14];      // Protocol-specific address data  
};
```

- Listen

The `listen()` function in socket programming is used to make the server node wait and listen for connections from the client node on the port and address specified by the `bind()` function. The syntax is,

```
int listen(int socket_descriptor, int back_log);
```

The **socket\_descriptor** represents the value of the file descriptor returned by the `socket()` function.



- *int listen(int socket\_descriptor, int backlog);*

The **backlog** marks the maximum number of connection requests that can be made to the server by client nodes at a time. The number of requests made after the number specified by backlog may cause an error or will be ignored by the server if the options for retransmission are set.

The function returns 0 on listening on the address and port specified or returns -1 on failure.

example

```
listen(server_fd, 3);
```

- Accept

The `accept()` function is used to establish a connection between the server and the client nodes for the transfer of data. This call typically blocks until a client connects with the server. The syntax is,

```
int accept(int socket_descriptor, struct sockaddr *address, socklen_t  
*length_of_address);
```

The **socket\_descriptor** represents the value of the file descriptor returned by the `socket()` function.

```
int accept(int socket_descriptor, struct sockaddr* address, socklen_t*  
length_of_address);
```

The **address** is the variable of the `sockaddr_in` structure in which the address of the socket returned from the function will be stored.

The **length\_of\_address** depicts the size of the address parameter.

The `accept()` function creates a new socket from the first connection request for the specified `socket_descriptor` and returns the file descriptor of the new socket. The file descriptor of this new socket is used in the `read()` and `write()` functions to send and receive data to and from the client node.

The `accept()` function is called by a server application after it has set up a listening socket using `listen()`. It extracts the first connection request from the queue of pending connections, creates a new socket for the connection, and returns a file descriptor for this new socket. The new socket can then be used for communication with the client.

example

```
struct sockaddr_in clientAddr;
```

```
socklen_t addrSize = sizeof(clientAddr);
```

```
nw_socket = accept(server_fd, (struct sockaddr_in*)&clientAddr, &addrSize);
```

- read

The `read()` function is used to receive data between client and server. The syntax of `read()` function is,

```
ssize_t read(int socket_descriptor, void *buffer, size_t size);
```

The **socket\_descriptor** represents the value of the socket descriptor returned by the `accept()` function.

The **buffer** represents the memory location where the data read is stored.

The **size** represents the maximum number of data bytes that can be stored in buffer.

The `read()` function, on success, returns the number of bytes read (zero indicates end of stream). It is not an error if this number is smaller than the number of bytes requested. On error, -1 is returned, and `errno` is set to indicate the error.

example

```
ssize_t valread;  
char buffer[1024];  
valread = read(nw_socket, buffer, 1024 - 1);  
// subtract 1 for the null terminator at the end
```

- write

The write() function is used to send data between client and server. The syntax of write() function is,

```
ssize_t write(int socket_descriptor, void *buffer, size_t count);
```

The **socket\_descriptor** represents the value of the socket descriptor returned by the accept() function.

The **buffer** represents the memory location where the data to be sent is stored.

The **count** represents the number of data bytes that are stored in buffer.

The `write()` function, on success, returns the number of bytes written. On error, the function `write()` returns -1, and `errno` is set to indicate the error.

example

```
char hello[] = "Hello from server";  
write(nw_socket, hello, strlen(hello));
```



- close

The `close()` function deallocates the socket descriptor passed as argument. To deallocate means to make the socket descriptor available for return by subsequent calls to `socket()`. The syntax is `close()` function is

```
int close(socket_descriptor);
```

The **socket\_descriptor** represents the value of the socket descriptor returned by the `socket()` or by the `accept()` function.

Upon successful completion, `close()` function returns 0; otherwise, -1 and `errno` is set to indicate the error.

- Socket Connection

Similar to the server-side, the client-side also needs to create a socket using the `socket()` function. This will create a socket that can send the connection request to the server. The client can connect the socket to the address of the server using the `connect()` system call.

```
// Create client socket  
cliSoc = socket (AF_INET, SOCK_STREAM, 0);  
if (cliSoc < 0) {  
    perror ("Error in socket creation");  
    exit (1);  
}
```

- Connect

The `connect()` function is used to send the connection request and connect to the server node. The syntax of the function is,

```
int connect(int socket_descriptor, const struct sockaddr *address, socklen_t  
length_of_address);
```

The **socket\_descriptor** represents the value of the file descriptor returned by the `socket()` function during the creation of a socket on the client-side.

The **address** represents the structure with the information of the address and port number of the server node to which the connection is to be made.

```
int connect(int socket_descriptor, const struct sockaddr *address, socklen_t  
length_of_address);
```

The **length\_of\_address** is the size of the address structure used in the second parameter.

```
// Set server address parameters  
serverAddr.sin_family = AF_INET;  
serverAddr.sin_port = htons (PORT);  
serverAddr.sin_addr.s_addr = inet_addr (SERVER_IP);  
// Connect to the server  
if (connect (cliSoc , (struct sockaddr*) & serverAddr, sizeof (serverAddr)) < 0) {  
    perror("Error in connecting to server");  
    exit (1);  
}
```

```
int connect(int socket_descriptor, const struct sockaddr *address, socklen_t  
length_of_address);
```

The connect() function returns a value of 0 on successfully connecting with the server and returns a value of -1 on error or the connection fails.

Similar to the server-side, the client-side also can invoke the read() and write() functions to send and receive data between client and server and the close() function to close the socket stream.

Let's see an implementation in which one hello message is exchanged between server and client to demonstrate the AF\_INET client/server model.

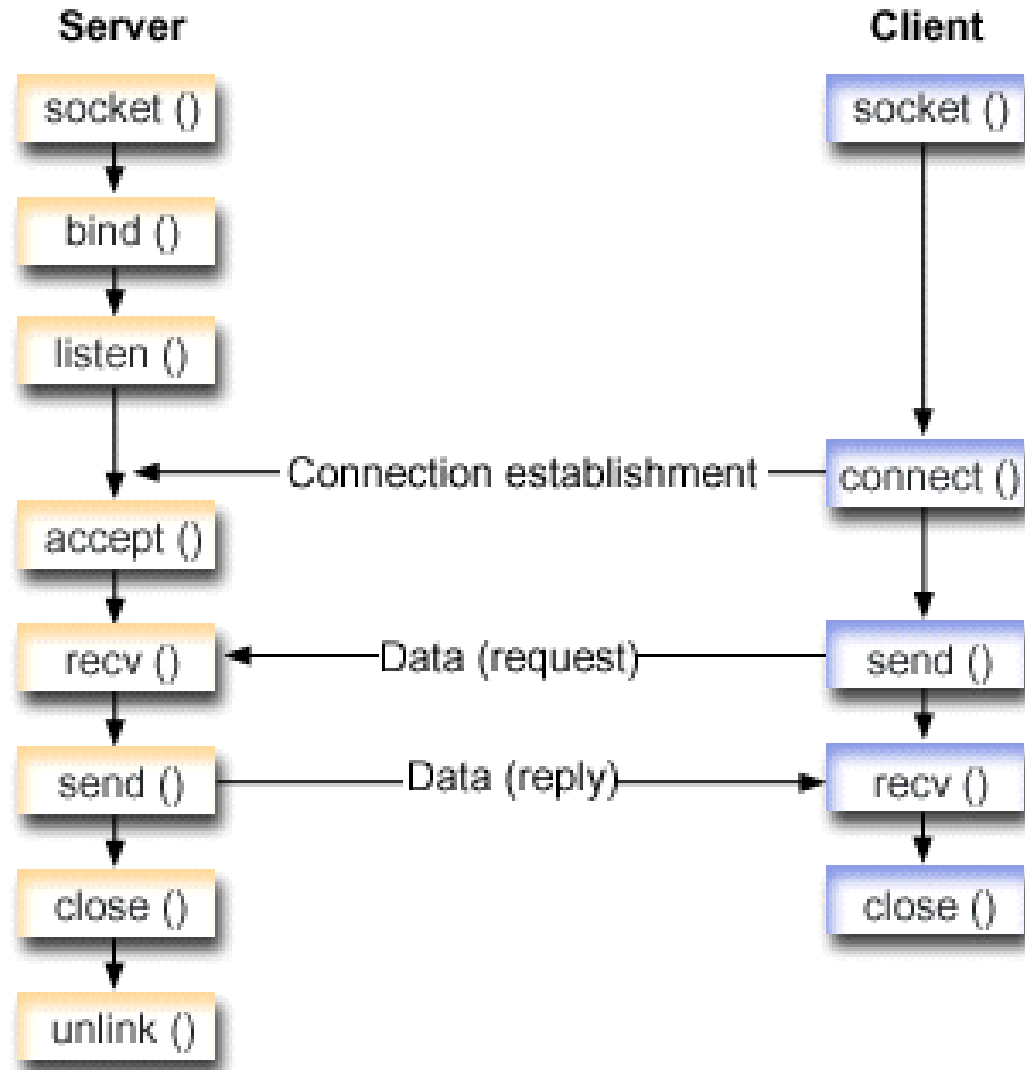
# Concurrent server stages



A concurrent server, as said, when processing the request, creates a child process (or a thread) responsible for providing the requested services, to immediately wait for further requests.

```
pid_t pid = fork();
if (pid == 0) {
    // Child process
    close (serverSocket);
    // receive messages from the client, read() and send() functions invocation
}
else if (pid > 0 ) {
    // Parent process
    close(clientSocket);
}
```

# Using AF\_UNIX address family



Sockets that use the AF\_UNIX address family can be connection-oriented (type SOCK\_STREAM) or connectionless (type SOCK\_DGRAM).

Both types are reliable because there are no external communication functions connecting the two processes.

Socket flow of events for a server application that uses AF\_UNIX address family.

The **socket()** API returns a socket descriptor, which represents an endpoint. The statement also identifies the UNIX address family with the stream transport (SOCK\_STREAM) being used for this socket. You can also use the **socketpair()** API to initialize a UNIX socket.

After the socket descriptor is created, the **bind()** API gets a unique name for the socket.



# Using AF\_UNIX address family



The name space for UNIX domain sockets consists of path names. When a sockets program calls the `bind()` API, an entry is created in the file system directory. If the path name already exists, the `bind()` fails. Thus, a UNIX domain socket program should always call an `unlink()` API to remove the directory entry when it ends.

The **`listen()`** allows the server to accept incoming client connections.

The server uses the **`accept()`** function to accept an incoming connection request. The `accept()` call will block indefinitely waiting for the incoming connection to arrive.

# Using AF\_UNIX address family



The **recv()** API receives data from the client application.

The **send()** API send data back to the client.

The **close()** API closes any open socket descriptors.

The **unlink()** API removes the UNIX path name from the file system.

Socket flow of events for a client application that uses AF\_UNIX address family

The **socket()** API returns a socket descriptor, which represents an endpoint. The statement also identifies the UNIX address family with the stream transport (SOCK\_STREAM) being used for this socket. You can also use the **socketpair()** API to initialize a UNIX socket.

After the socket descriptor is received, the **connect()** API is used to establish a connection to the server.

The **send()** API sends data bytes to the server.

# Using AF\_UNIX address family



The **recv()** API receives data bytes back from the server.

The **close()** API closes any open socket descriptors.

Let's see an implementation in which one hello message is exchanged between server and client to demonstrate the AF\_UNIX client/server model.

# Some useful functions



```
#include <sys/socket.h>
```

```
#include <netdb.h>
```

```
int getaddrinfo(const char *nodename, const char *servname,  
    const struct addrinfo *hints, struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *ai);
```

```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

The **getaddrinfo()** function translates the name of a service location (for example, a host name) and/or a service name and returns a set of socket addresses and associated information to be used in creating a socket with which to address the specified service.

The `nodename` and `servname` arguments are either null pointers or pointers to null-terminated strings. One or both of these two arguments must be a non-null pointer.

```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

It is part of the modern POSIX standard and is preferred over older functions like `gethostbyname()` and `getservbyname()` because it is protocol-independent (works for both IPv4 and IPv6) and thread-safe.

If the ***nodename*** argument is not null, it can be a descriptive name or can be an address string. Address strings using internet standard dot notation are valid if the specified address family is `AF_INET` or `AF_UNSPEC`.

If `nodename` is not null, the requested service location is named by `nodename`; otherwise, the requested service location is local to the caller.

nodename:

- The hostname (e.g., "www.example.com") or IP address (e.g., "192.168.1.1" for IPv4 or "::1" for IPv6).
- If set to NULL, it means the local host.



```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

If ***servname*** argument is null, the call returns network-level addresses for the specified ***nodename***. If ***servname*** is not null, it is a null-terminated character string identifying the requested service. This can be either a descriptive name or a numeric representation suitable for use with the address family or families. If the specified address family is `AF_INET`, `AF_INET6` or `AF_UNSPEC`, the service can be specified as a string specifying a decimal port number.

- The service name (e.g., "http", "ftp") or the port number as a string (e.g., "80").
- If `NULL`, no service is specified.

# Some useful functions



```
int getaddrinfo(const char *nodename, const char *servname,  
               const struct addrinfo *hints, struct addrinfo **res);
```

The `addrinfo` structure used by `getaddrinfo()` contains the following fields:

```
struct addrinfo {  
    int      ai_flags;  
    int      ai_family;  
    int      ai_socktype;  
    int      ai_protocol;  
    socklen_t ai_addrlen;  
    struct sockaddr *ai_addr;  
    char      *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

If the argument ***hints*** is not null, it refers to a structure containing input values that may direct the operation by providing options and by limiting the returned information to a specific socket type, address family and/or protocol. In this hints structure every member other than `ai_flags`, `ai_family`, `ai_socktype` and `ai_protocol` must be zero or a null pointer.

```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

## *ai\_family*

This field specifies the desired address family for the returned addresses. Valid values for this field include `AF_INET` and `AF_INET6`. The value `AF_UNSPEC` indicates that `getaddrinfo()` should return socket addresses for any address family (either IPv4 or IPv6, for example) that can be used with node and service.

# Some useful functions



```
int getaddrinfo(const char *nodename, const char *servname,  
               const struct addrinfo *hints, struct addrinfo **res);
```

## *ai\_socktype*

This field specifies the preferred socket type, for example SOCK\_STREAM or SOCK\_DGRAM. Specifying 0 in this field indicates that socket addresses of any type can be returned by getaddrinfo().

## *ai\_protocol*

This field specifies the protocol for the returned socket addresses. Specifying 0 in this field indicates that socket addresses with any protocol can be returned by getaddrinfo().

# Some useful functions



```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

If hints is a null pointer, the behavior must be as if it referred to a structure containing the value zero for the ai\_flags, ai\_socktype and ai\_protocol fields, and AF\_UNSPEC for the ai\_family field.

```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

The `getaddrinfo()` function allocates and initializes a linked list of `addrinfo` structures, one for each network address that matches node and service, subject to any restrictions imposed by hints, and returns a pointer to the start of the list in ***res***. The items in the linked list are linked by the `ai_next` field.

There are several reasons why the linked list may have more than one `addrinfo` structure. Normally, the application should try using the addresses in the order in which they are returned.

# Some useful functions



```
int getaddrinfo(const char *nodename, const char *servname,  
               const struct addrinfo *hints, struct addrinfo **res);
```

If *hints.ai\_flags* includes the `AI_CANONNAME` flag, then the `ai_canonname` field of the first of the `addrinfo` structures in the returned list is set to point to the official name of the host.

The *ai\_family*, *ai\_socktype*, and *ai\_protocol* fields return the socket creation parameters (i.e., these fields have the same meaning as the corresponding arguments of `socket()` function). For example, `ai_family` might return `AF_INET` or `AF_INET6`; `ai_socktype` might return `SOCK_DGRAM` or `SOCK_STREAM`; and `ai_protocol` returns the protocol for the socket.



# Some useful functions



```
int getaddrinfo(const char *nodename, const char *servname,  
const struct addrinfo *hints, struct addrinfo **res);
```

A pointer to the socket address is placed in the *ai\_addr* field, and the length of the socket address, in bytes, is placed in the *ai\_addrlen* field.

`getaddrinfo()` returns 0 if it succeeds, or one of the following nonzero error codes: `EAI_ADDRFAMILY`, `EAI_AGAIN`, `EAI_BADFLAGS`, `EAI_FAIL`, `EAI_FAMILY`, `EAI_MEMORY`, `EAI_NODATA`, `EAI_NONAME`, `EAI_SERVICE`, `EAI_SOCKTYPE`, `EAI_SYSTEM`.

The `gai_strerror()` function translates these error codes to a human readable string, suitable for error reporting.

# Some useful functions



```
int getaddrinfo(const char *nodename, const char *servname,  
               const struct addrinfo *hints, struct addrinfo **res);  
  
struct addrinfo *result = NULL, *ptr = NULL, hints;  
...  
hints.ai_family = AF_UNSPEC;  
hints.ai_socktype = SOCK_STREAM;  
hints.ai_protocol = IPPROTO_TCP;  
rc = getaddrinfo("myhost.mydomain.com", "8080", &hints, &result);  
for (ptr = result; ptr != NULL; ptr = ptr->ai_next) {  
    mySocket = socket(ptr->ai_family, ptr->ai_socktype, ptr->ai_protocol);  
    ...  
}
```

```
#include <sys/socket.h>
```

```
#include <netdb.h>
```

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
    char *node, socklen_t nodelen, char *service,  
    socklen_t servicelen, unsigned int flags);
```

The **getnameinfo()** function is the inverse of `getaddrinfo()` function. It translates a socket address to a node name and service location, all of which are defined as with `getaddrinfo()`.

The argument `sa` points to a socket address structure to be translated.

# Some useful functions



```
int main(int argc, char *argv[]) {
    struct addrinfo hints, *res, *p;
    int status;
    char ipstr[INET6_ADDRSTRLEN]; // Buffer to hold IP address (IPv4 or IPv6)

    if (argc != 2) {
        fprintf(stderr, "Usage: %s hostname\n", argv[0]);
        return 1;
    }

    // Initialize the hints structure
    memset(&hints, 0, sizeof hints);
    hints.ai_family = AF_UNSPEC; // Allow IPv4 or IPv6
    hints.ai_socktype = SOCK_STREAM; // TCP socket

    // Call getaddrinfo to resolve hostname
    if ((status = getaddrinfo(argv[1], "http", &hints, &res)) != 0) {
        fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(status));
        return 2;
    }

    printf("IP addresses for %s:\n\n", argv[1]);
```

```
// Loop through all the results
for (p = res; p != NULL; p = p->ai_next) {
    void *addr;
    char *ipver;
    // Get the pointer to the address itself
    if (p->ai_family == AF_INET) { // IPv4
        struct sockaddr_in *ipv4 = (struct sockaddr_in *)p->ai_addr;
        addr = &(ipv4->sin_addr);
        ipver = "IPv4";
    } else { // IPv6
        struct sockaddr_in6 *ipv6 = (struct sockaddr_in6 *)p->ai_addr;
        addr = &(ipv6->sin6_addr);
        ipver = "IPv6";
    }

    // Convert the IP to a string and print it
    inet_ntop(p->ai_family, addr, ipstr, sizeof ipstr);
    printf(" %s: %s\n", ipver, ipstr);
}
freeaddrinfo(res); // Free the linked list
return 0;
}
```

# Some useful functions



```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
char *node, socklen_t nodelen, char *service,  
socklen_t servicelen, unsigned int flags);
```

The `getnameinfo()` function in socket programming is used to convert a socket address (e.g., an IP address and port) back into a human-readable string such as a hostname and service name (or port number). It essentially performs the reverse operation of `getaddrinfo()`.

It translates a binary socket address (e.g., `struct sockaddr_in` for IPv4 or `struct sockaddr_in6` for IPv6) into a readable format, replacing older functions like `gethostbyaddr()` and `inet_ntoa()` and working seamlessly with IPv4 and IPv6.

# Some useful functions



```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
char *node, socklen_t nodelen, char *service,  
socklen_t servicelen, unsigned int flags);
```

- sa (Socket Address): a pointer to the socket address structure (e.g., struct sockaddr\_in or struct sockaddr\_in6).
- salen: the size (in bytes) of the socket address structure.
- node: a pointer to a buffer where the hostname will be stored.
- nodelen: the size of the buffer allocated for the hostname.
- service: a pointer to a buffer where the service name (e.g., "http") or port number (e.g., "80") will be stored.
- servicelen: the size of the buffer allocated for the service name.
- flags: a set of flags that control the behavior of the function (explained below).

# Some useful functions



```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
char *node, socklen_t nodelen, char *service,  
socklen_t servicelen, unsigned int flags);
```

If the argument `node` is non-NULL and the argument `nodelen` is nonzero, then the argument `node` points to a buffer able to contain up to `nodelen` characters that will receive the node name as a null-terminated string. If the argument `node` is NULL or the argument `nodelen` is zero, the node name will not be returned. If the node's name cannot be located, the numeric form of the node's address is returned instead of its name.

# Some useful functions



```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
char *node, socklen_t nodelen, char *service,  
socklen_t servicelen, unsigned int flags);
```

If the argument `service` is non-NULL and the argument `servicelen` is nonzero, then the argument `service` points to a buffer able to contain up to `servicelen` characters that will receive the service name as a null-terminated string. If the argument `service` is NULL or the argument `servicelen` is zero, the service name will not be returned. If the service's name cannot be located, the numeric form of the service address (for example, its port number) is returned instead of its name.

The arguments `node` and `service` cannot both be NULL.



```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
char *node, socklen_t nodelen, char *service,  
socklen_t servicelen, unsigned int flags);
```

The flags argument is a flag that changes the default actions of the function. By default the fully-qualified domain name (FQDN) for the host is returned, but

- If the flag bit NI\_NOFQDN is set, only the nodename portion of the FQDN is returned for local hosts.
- If the flag bit NI\_NUMERICHOST is set, the numeric form of the host's address is returned instead of its name, under all circumstances.

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
    char *node, socklen_t nodelen, char *service,  
    socklen_t servicelen, unsigned int flags);
```

- If the flag bit NI\_NAMEREQD is set, an error is returned if the host's name cannot be located.
- If the flag bit NI\_NUMERICSERV is set, the numeric form of the service address is returned (for example, its port number) instead of its name, under all circumstances.
- If the flag bit NI\_DGRAM is set, this indicates that the service is a datagram service (SOCK\_DGRAM). The default behavior is to assume that the service is a stream service (SOCK\_STREAM).

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
char *node, socklen_t nodelen, char *service,  
socklen_t servicelen, unsigned int flags);
```

On success, 0 is returned, and node and service names, if requested, are filled with null-terminated strings, possibly truncated to fit the specified buffer lengths. On error, one of the following nonzero error codes is returned:

EAI\_AGAIN, EAI\_BADFLAGS, EAI\_FAIL, EAI\_FAMILY, EAI\_MEMORY,  
EAI\_NONAME, EAI\_OVERFLOW, EAI\_SYSTEM.

The `gai_strerror()` function translates these error codes to a human readable string, suitable for error reporting.